

GeekBand 极客班

互联网人才加油站!



C++系统工程师



iOS开发工程师



Android开发工程师



PM产品经理

系统设计中七剑客

GeekBand

极客班

大纲

同步

网络

数据库

分布式

性能

估算

面向对象

案例

- 社交网站信息流
- 日志统计
- 网络爬虫
- 电商产品页面

Introduction

System design:

1-3 rounds in interviews

For new-grad: not required/simple

For experienced: important to show your knowledge and thoughtful ideas

Knowledge, Design, Communication!

Concurrency

Thread vs. Process

Consumer and Producer

Blockingqueue

Tracking:


Synchronized, Asynchronized

GeekPanda

极客班

Network

The Seven Layers of OSI



Visit URL

What happens after you typed a URL in your browser and pressed return key?

GeekBand

极客班

Database

Relational DB vs. KV Store

Sharding vs. Clustering

TinyURL:

Store the mapping from shortlink code to full URL. The record/
document:

code: varchar(8)

url: varchar(1000)

created_at: timestamp

We also need to store the reverse mapping from URL back to code.

Distribute System

How to scale Tiny URL service?

- Stateless frontend servers behind a load balancer

- Sharded/replicated database (on shortlink code)

- Memcached to scale read traffic

- Spread write load

- Locally buffered event tracking + async flush to high-throughput message queue

- Use a distributed unique ID generator (64-bit)

Performance

Numbers Everyone Should Know

Cache is KEY!

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zip	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Estimation

How many piano tuners are there in the entire world?

Tiny URL: How much is total storage?

URL Length 10 - 1000 chars.

Total accumulated URL number 100 M

New URL registrations are on the order 100,000/day (1/sec)

Redirect requests are on the order of 100M/day (1000/sec)

Design Pattern

23 patterns:

MVC

Singleton

Factory

Iterator

Decorator

Facade

GeekBand

极客班

案例

News Feeds

Stats Server

Web Crawler

Amazon Product Page

GeekBand

极客班

News feed

Define feed

Organize

- aggregate

- dedup

- sort

GeekBand

极客班

Level 1.0

Database Schema:

User

Friendship

News

Get Newsfeed:

merge news

Newsfeed vs News

<i>UserId</i>	<i>Name</i>	<i>Age</i>
1	Jason	25
2	Michael	26

<i>FriendshipId</i>	<i>SourceId</i>	<i>TargetId</i>
1	1	2
2	2	1

<i>NewsId</i>	<i>AuthorId</i>	<i>Content</i>
1	2	"Hehe"
2	1	"Lala"

Why bad?

100+ friends

1 Query --> Get friends list

1 Query -->

SELECT news

WHERE timestamp>xxx

AND sourceid IN friend list

LIMIT 1000

IN is slow

Either Sequential scan or 100+ index queries

Level 2.0

Pull vs Push

Pull: Get news from each friend, merge them together.

(NewsFeed generated when user request)

Push: NewsFeed generated when news generated. (we have another table to store newsfeed, may cause duplicate news)

Push:

1 Query to select latest 1000 newsfeed.

100+ insert queries (Async)

Disadvantage: News Delay.

Level 3.0

Popular star (Justin Bieber)

Flowers 13M +

Async Push may cause over 30 minutes (13M+ insertions, delay too long)

Push + Pull

for popular star, don't push news to flowers

for every newsfeed request, merge non-popular user newsfeed (push) and popular users newsfeed (pull)

Level 4.0

Push disadvantage:

- Realtime

- Storage (Duplicate)

- Edit

Go back to PULL:

- Cache users' latest (14 days) news

- Broadcast multiple request to multiple servers (Shard by userId).

- Merge & sort newsfeed

- Cache newsfeeds for this user with timestamp

Click Stats Server

How are click stats stored?

A poor candidate will suggest write-back to a data store on every click

A good candidate will suggest some form of aggregation tier that accepts clickstream data, aggregates it, and writes back a persistent data store periodically

A great candidate will suggest a low-latency messaging system to buffer the click data and transfer it to the aggregation tier.

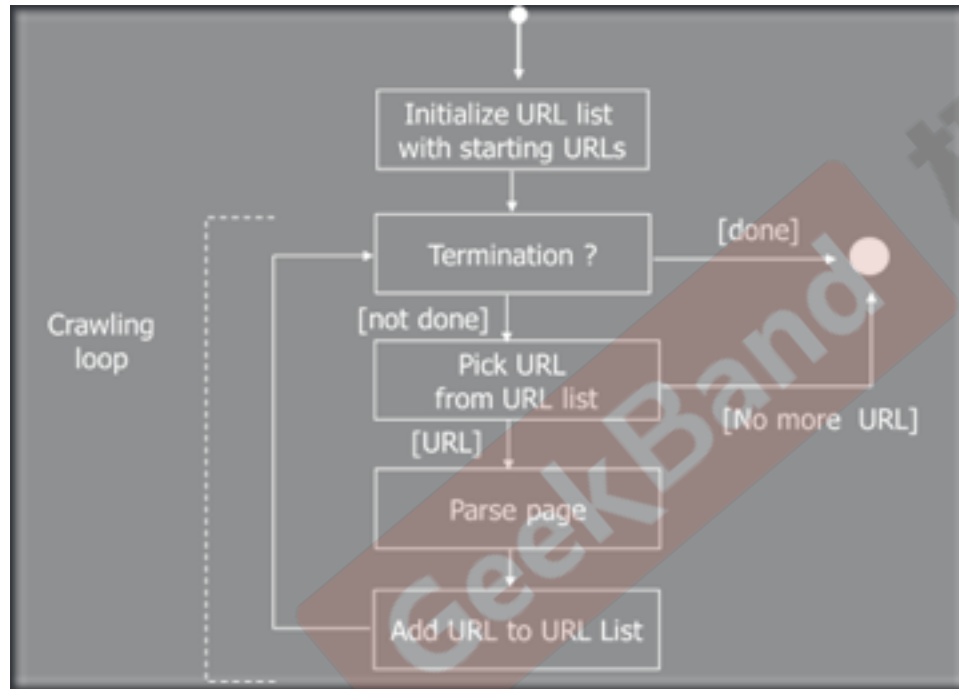
If daily, storing in hdfs and running map/reduce jobs to compute stats is a reasonable approach

If near real-time, the aggregation logic should compute stats

Cache Requirement

- [a] When a request comes look it up in the cache and if it hits then return the response from here and do not pass the request to the system
- [b] If the request is not found in the cache then pass it on to the system
- [c] Since cache can only store the last n requests, Insert the $n+1$ th request in the cache and delete one of the older requests from the cache
- [d] Design one cache such that all operations can be done in $O(1)$ – lookup, delete and insert.

Web Crawler



Amazon Product Page

The product page includes information such as

- a) product information
- b) user information
- c) recommended products (what do other customers buy after viewing this item, recommendations for you like this product, etc)

Reference

<http://highscalability.com/>

[The Log: What every software engineer should know about real-time data's unifying abstraction](#)

[Job Interviews: How should I prepare system design questions for Google/Facebook Interview?](#)

[HOW TO ACE A SYSTEMS DESIGN INTERVIEW](#)

[<Design Pattern>](#)

[<Design_Patterns_For_Dummies.pdf>](#)

<http://www.hiredintech.com/app>